

# **SCNN: An Accelerator for Compressed-sparse Convolutional Neural Networks**

Original Paper by: Angshuman Parashar et al. (NVIDIA, ISCA '17)

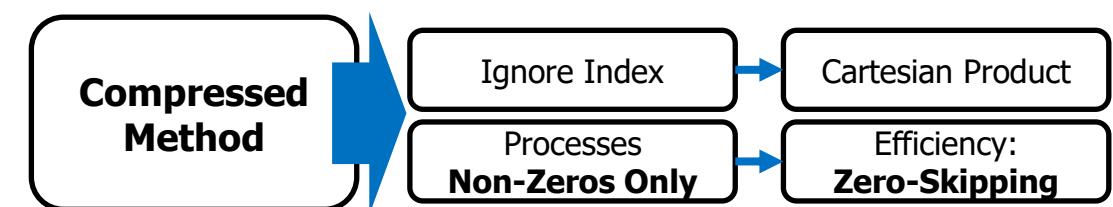
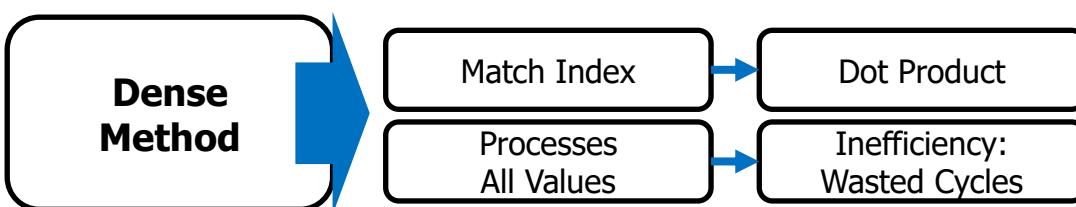
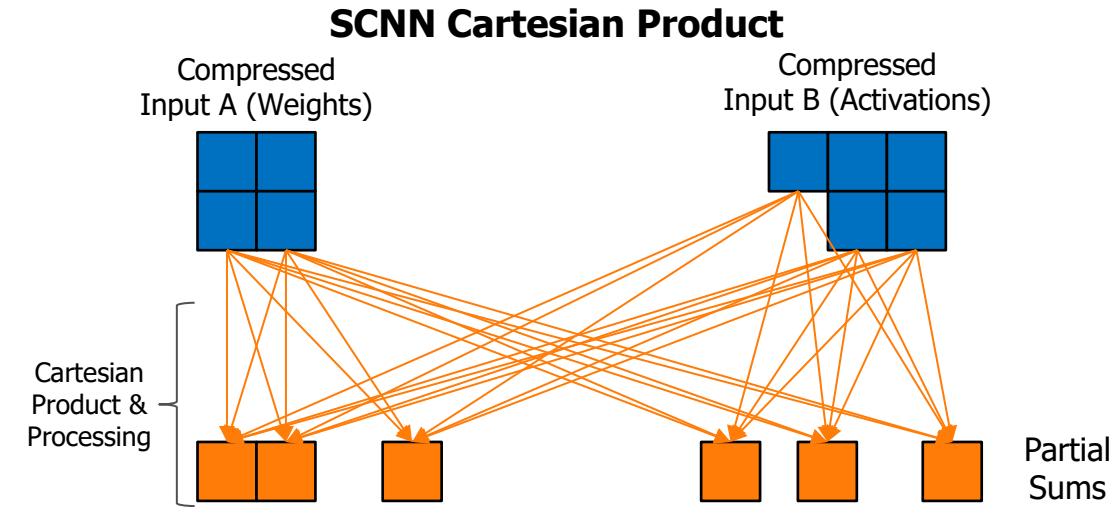
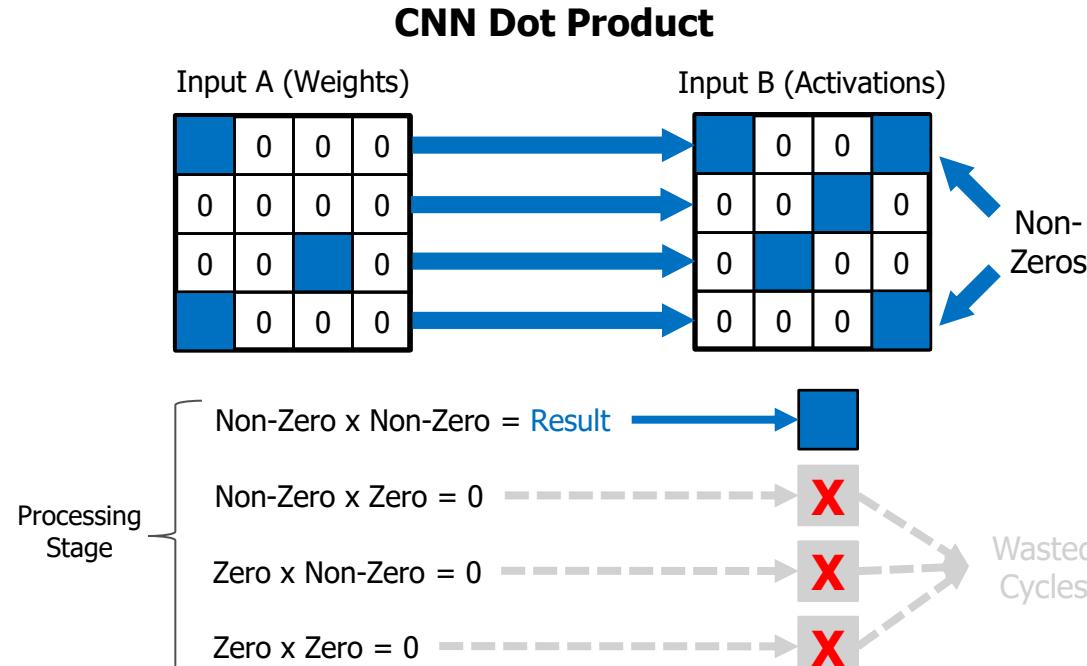
Follow-up Presentation

Presented by Jongyun Hur

# **Cartesian Product & Zero-Skipping**

# Cartesian Product & Zero-Skipping

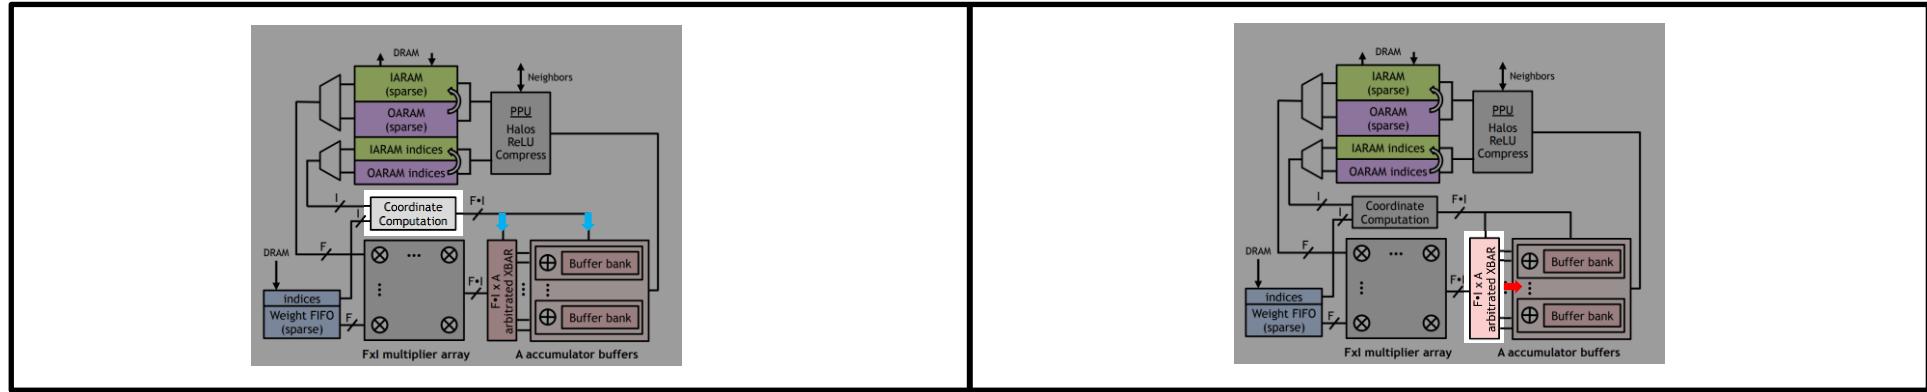
Why cartesian product enables Zero-Skipping



# **Coordinate Computation vs. XBAR**

# Coordinate Computation vs. XBAR

Decoupling address calculation from data delivery



## Coordinate Computation

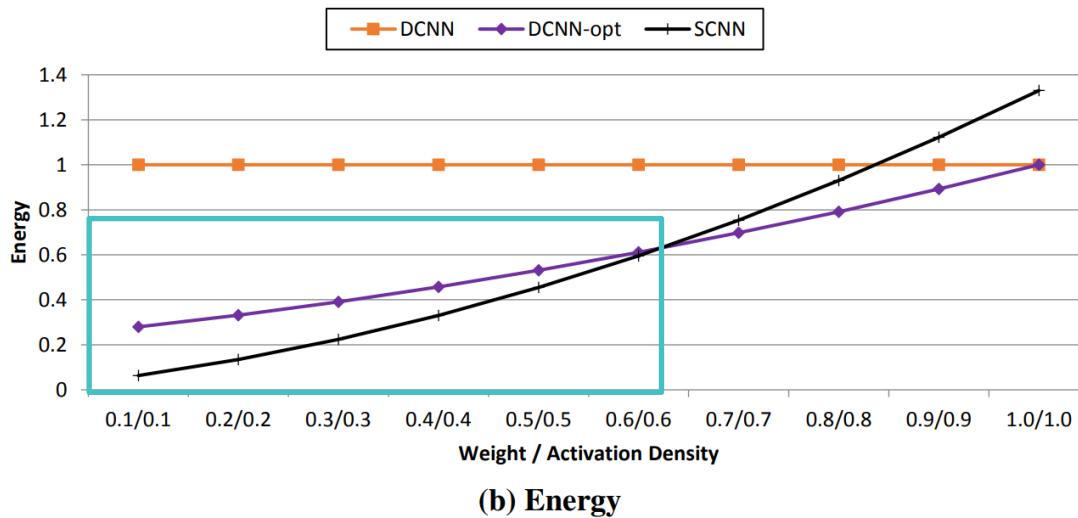
## XBAR

	Coordinate Computation	XBAR
Primary Role	<b>Computes</b> where the data should go.	<b>Delivers</b> the actual value to the destination.
Destinations & Sent Info	Sends Distinct Info to <b>Two</b> Targets: 1. To XBAR: Routing control signal (Bank ID) 2. To Buffer: Storage location (Memory Address)	Sends Data to Single Target: 1. To Buffer: Actual partial sum values generated by multipliers.
Mechanism	Decodes compressed indices to generate output coordinates (w,h,k) and splits them into routing ID and local address.	Routes the partial sums to the specific accumulator bank designated by the routing ID.

# **DCNN-opt vs. SCNN**

# DCNN-opt vs. SCNN

Why DCNN-opt consumes more energy than SCNN

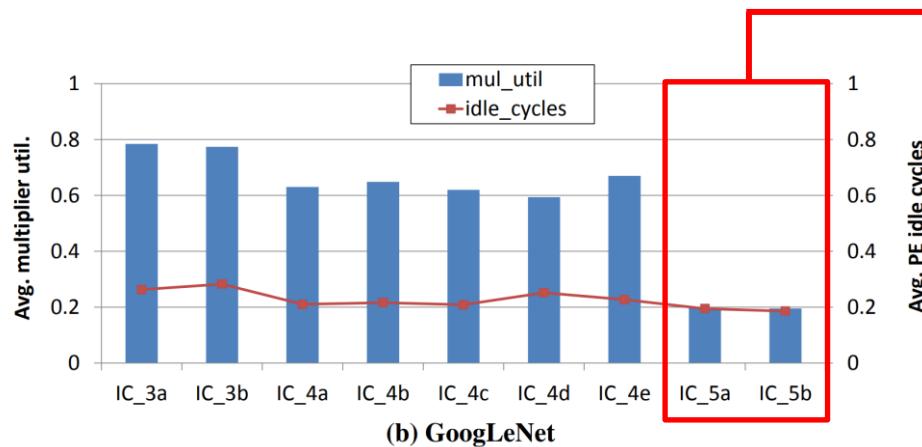


- **DCNN-opt (Zero-Gating)**
  - Reduces dynamic energy via Zero-Gating compared to DCNN.
  - Cannot skip zero operations → Execution time remains unchanged.
  - **Result: Wastes static power during idle cycles.**
- **SCNN (Zero-Skipping)**
  - Eliminates dead cycles via Zero-Skipping.
  - **Reduces both energy consumption and execution time.**
  - Result: Achieves the lowest total energy use by minimizing static power overhead.

# GoogLeNet & Impact of Filter Size

# GoogLeNet & Impact of Filter Size

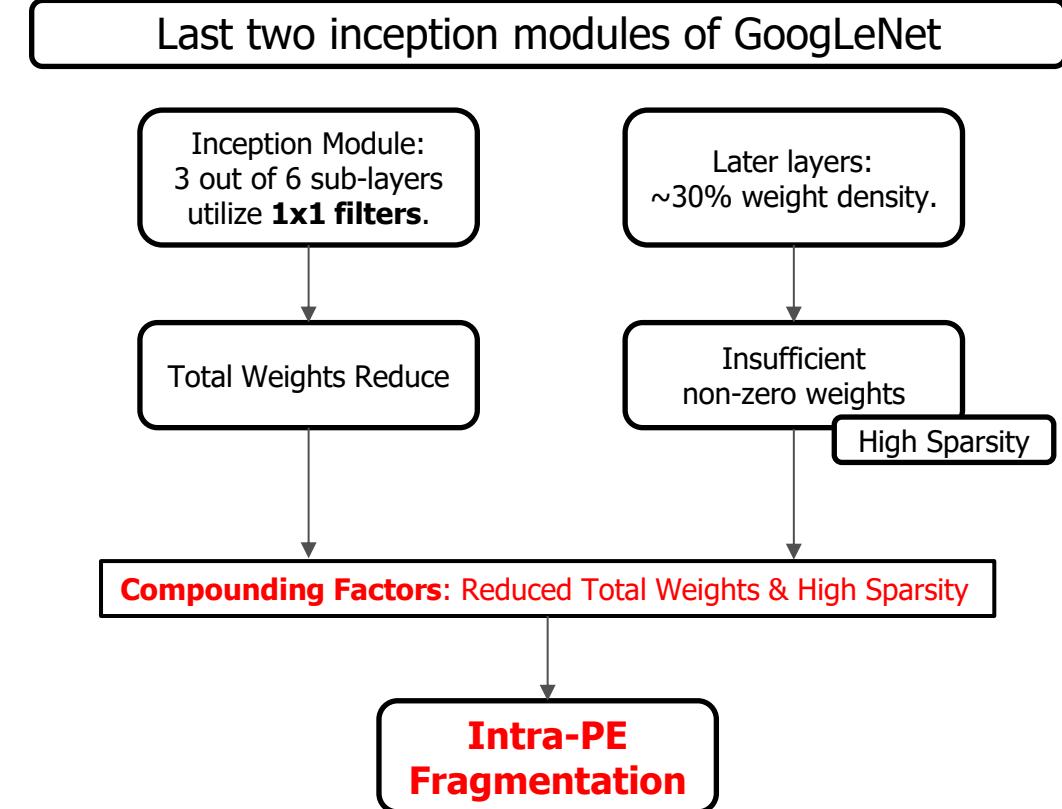
How small filters & high sparsity induce fragmentation



(b) GoogLeNet

$$\begin{aligned} \text{Total Weights} &= \text{Filter Width } (R) \times \text{Filter Height } (S) \times \text{Group Size } (K_c) \\ &= \text{Filter Width } (R) \times \text{Filter Height } (S) \times 8 \end{aligned}$$

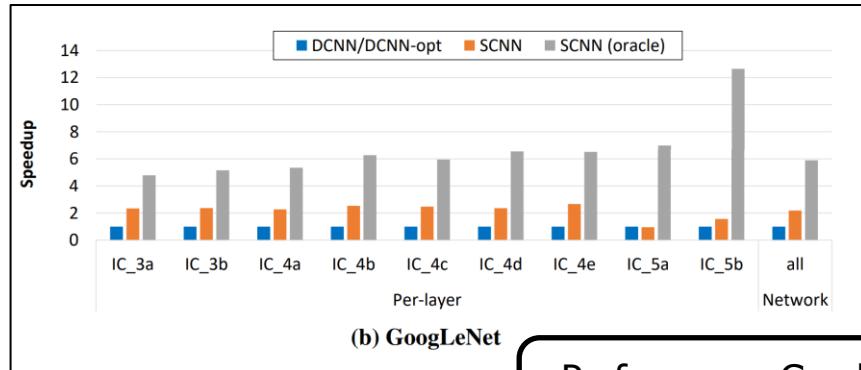
- Group size ( $K_c$ ) is fixed at 8 in the paper.
- Total Weights: Total learnable parameters per layer



# **SCNN vs. Oracle**

# SCNN vs. Oracle

Performance gaps relative to the theoretical limit



Performance Gap between SCNN and SCNN (oracle)

$$\text{SCNN(oracle)} = \frac{\text{Total Multiplications}}{\text{Number of Multipliers}} = \frac{\text{Total Multiplications}}{1024}$$

An ideal theoretical abstraction  
without physical hardware limitations

## SCNN

**Distributed** Resources:  
16 Multipliers per PE

Leads to: Low Utilization

**Uneven** Workload Distribution  
across PEs

Leads to: Idle Cycles

Intra-PE fragmentation

Inter-PE load imbalance

## SCNN (oracle)

**Unified** Resources:  
1024 Multipliers in Single Pool

No Intra-PE fragmentation

**Even** Workload Distribution  
across PEs

No Inter-PE load imbalance

# **Memory Difference: CNN vs. SCNN**

# Memory Difference: CNN vs. SCNN

Functional specialization of L0, L1 storage and accumulation logic

	<b>CNN</b>	<b>SCNN</b>
<b>L1 Memory</b>	Global Buffer	IARAM/OARAM (Physically distributed in PEs, but functionally serves as L1)
<b>L0 Memory</b>	Register File (RF)	Weight FIFO, Accumulator Buffer, Vector I/F Register
<b>Architecture</b>	<b>Unified</b> All data types (Input, Weight, Psum) share the same RF.	<b>Distributed</b> Uses physically separated, dedicated memories optimized for each data role.
<b>Accumulation Strategy</b>	<b>Deterministic (Fixed).</b> Output addresses are pre-determined, allowing sequential accumulation in the RF/Buffer.	<b>Scatter Accumulation.</b> Output addresses are scattered, requiring Banked Buffers and a Crossbar to handle conflicts.

