

# **In-Datacenter Performance Analysis of a Tensor Processing Unit**

Original Paper by: Norman P. Jouppi et al. (Google, ISCA '2017)

Presented by Jongyun Hur

# Introduction

# Introduction

## Why build the TPU?

**2006: Rejected Custom HW**

**FPGAs/ASICs Considered**  
(Only niche apps needed custom HW)

**2013: TPU Project Started**

**DNN Accuracy Explosion**  
(Voice Search: +3 min/day?)

**2015: Datacenter Deployment**

**Excess Capacity**  
(Abundant idle CPU resources)

**"Free" Compute**  
(Spare cycles cost almost nothing)

**Conclusion: CPU Suffices**  
(No custom HW needed)

**2016: Data Collection**

**Compute Demand Crisis**  
(Datacenter demand would **Double**)

**CPU Scaling Limits**  
(Cost prohibitive to scale with CPUs)

**2017: Paper Publication**

**Solution: Build TPU**  
(10x Perf/Watt vs GPU,  
15-month turnaround)

# Introduction

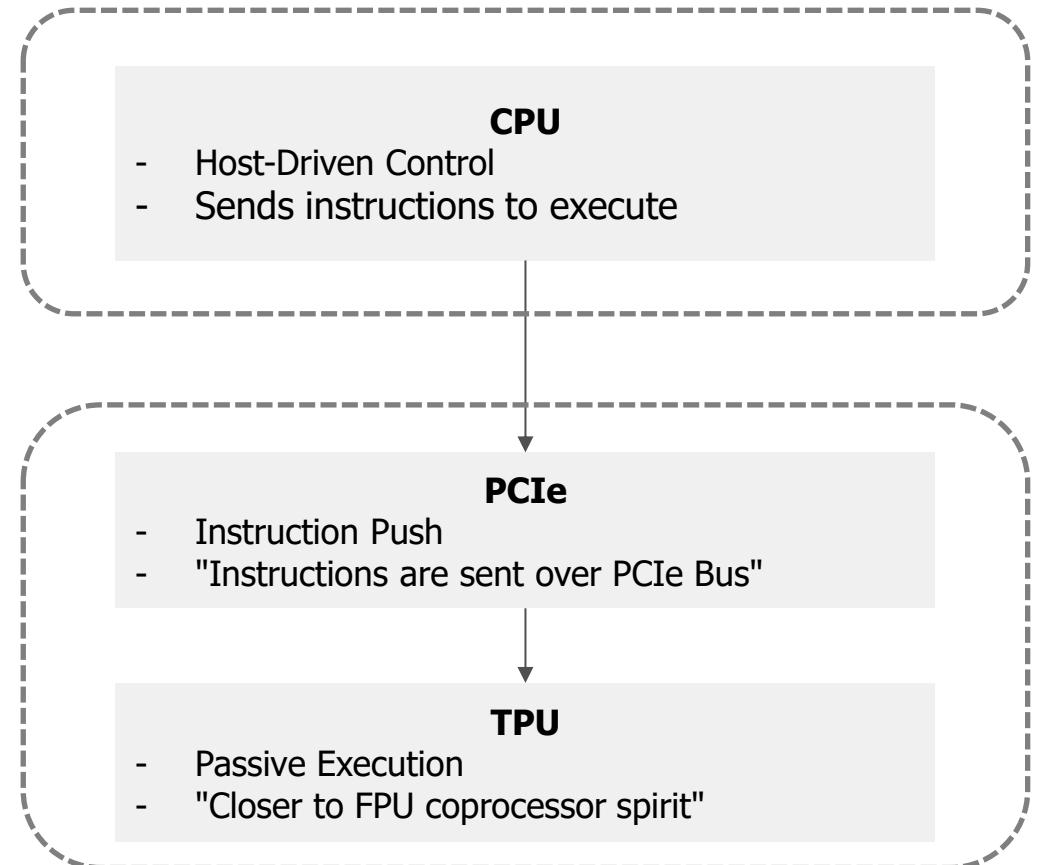
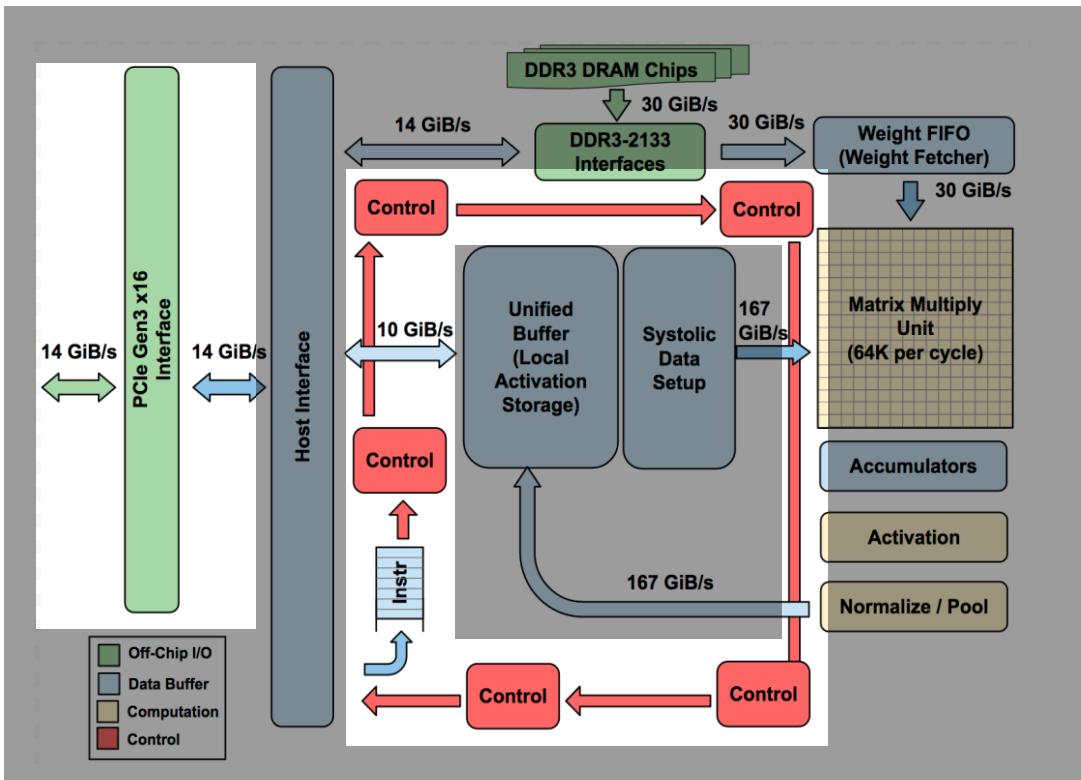
The 15-month mandate: Design trade-offs

Design Area	The Trade-off Decision	Reason for Speed
1. Connectivity	CPU Integration → <b>PCIe Coprocessor</b>	To plug into existing servers and avoid delaying deployment.
2. Control	Self-fetching Instructions → <b>Host-Driven (FPU-like)</b>	To <b>simplify hardware design</b> and debugging.
3. Memory	High Bandwidth (GDDR5) → <b>Conservative DDR3</b>	To use a proven interface and minimize circuit design risks.
4. Features	Power Gating → <b>Features Omitted</b>	Explicitly omitted to meet the strict " <b>time-to-deployment</b> " goal.

# Architecture

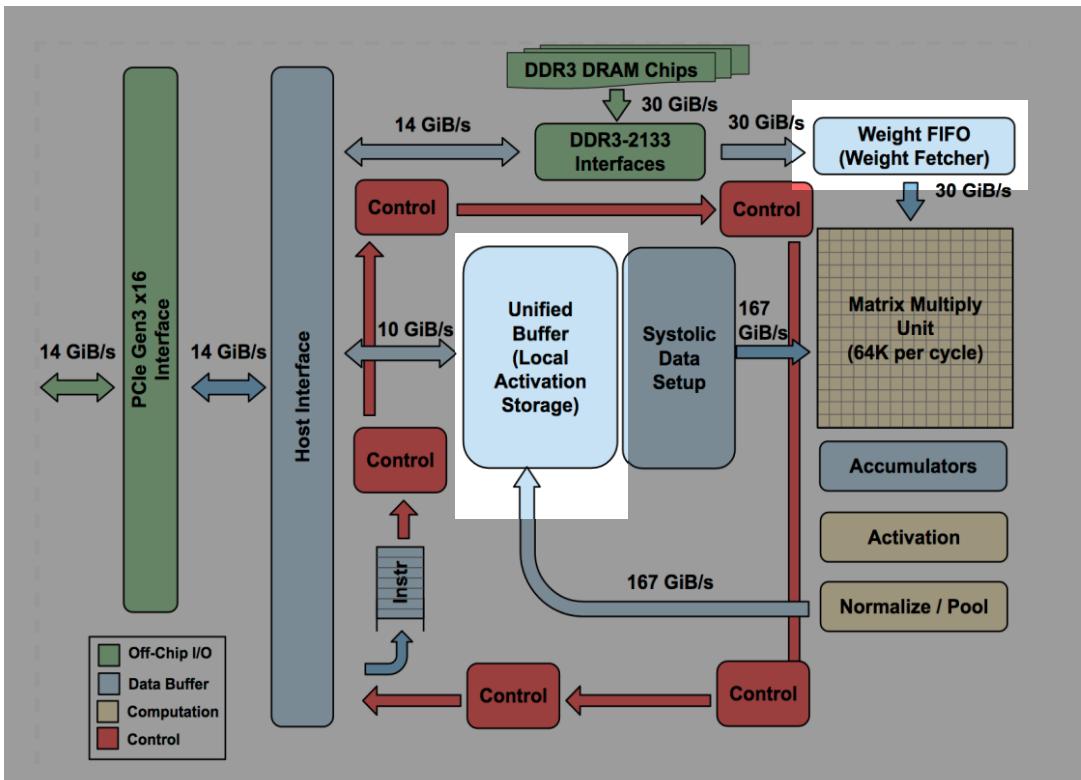
# Architecture

## TPU: System control & Instruction flow



# Architecture

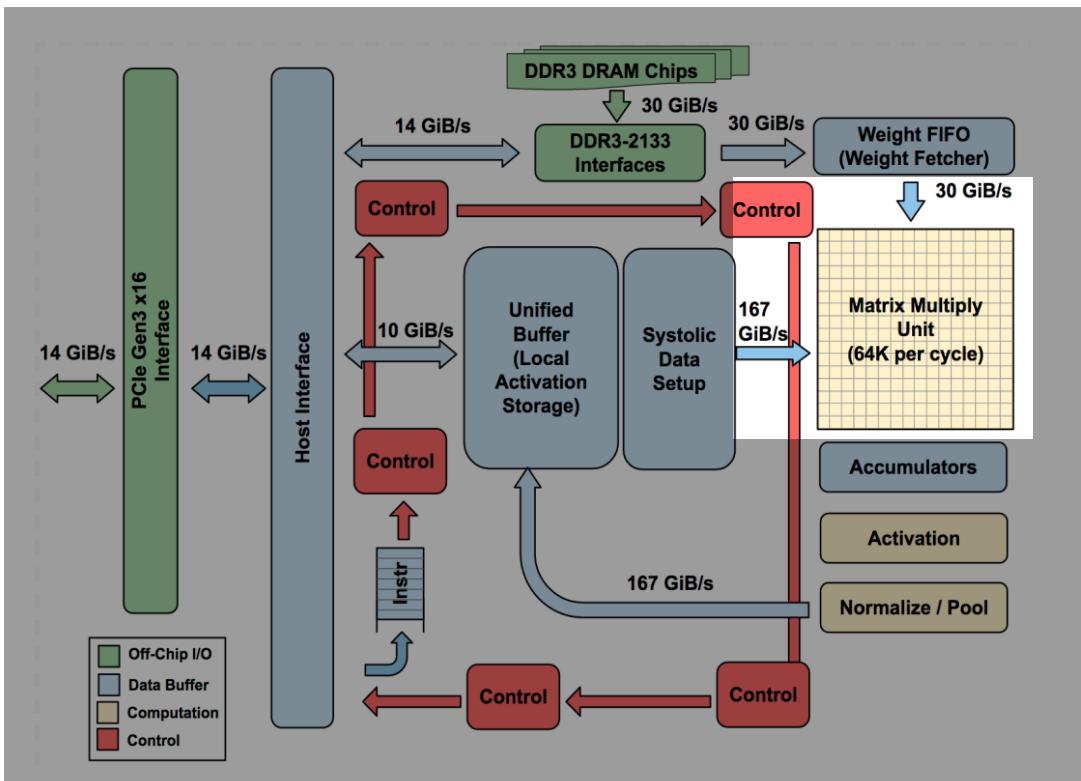
## TPU: Unified Buffer & Weight FIFO



Category	Unified Buffer	Weight FIFO
Data Type	Activations	Weights
Source	Host CPU (via PCIe Bus)	Off-chip DRAM (Weight Memory)
Key Feature	Software-Managed	Prefetching (Hides Latency)

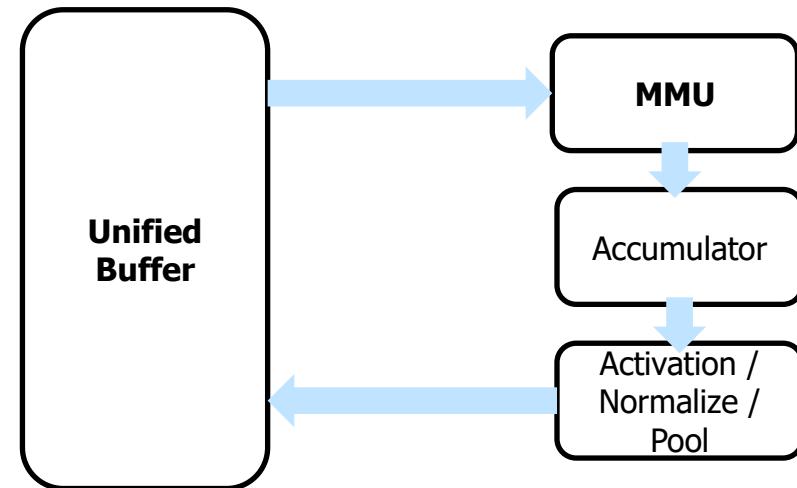
# Architecture

## TPU: Deterministic matrix multiplication



### Matrix Multiply Unit (MMU)

- **65,536 MACs** ( $256 \times 256$  Array)
- **Systolic Data Flow**
- **8-bit Integer Math**  
(Optimization for Inference)



# Architecture

## Matrix Multiply Unit: 256x256 Systolic Array

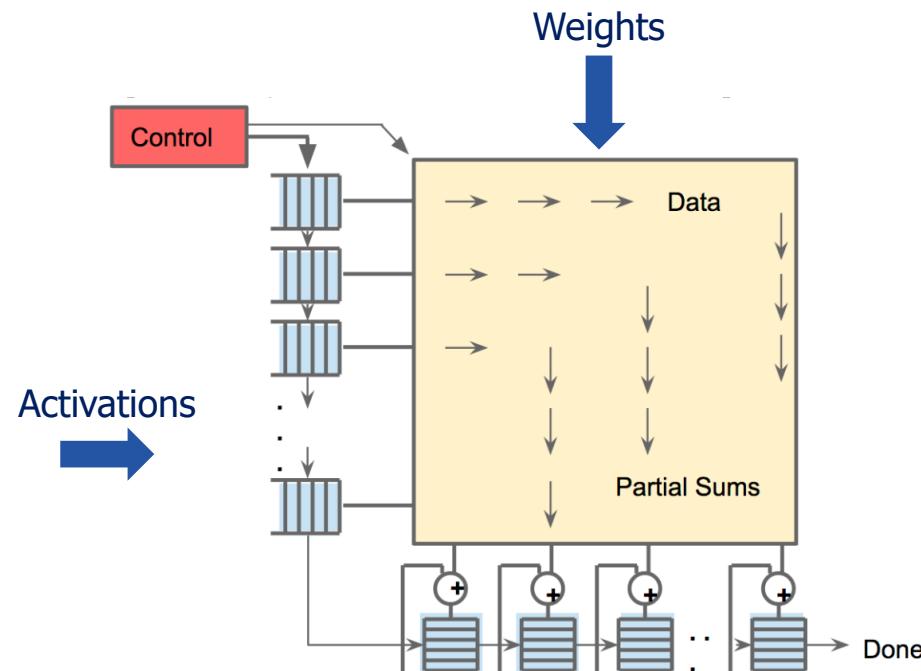


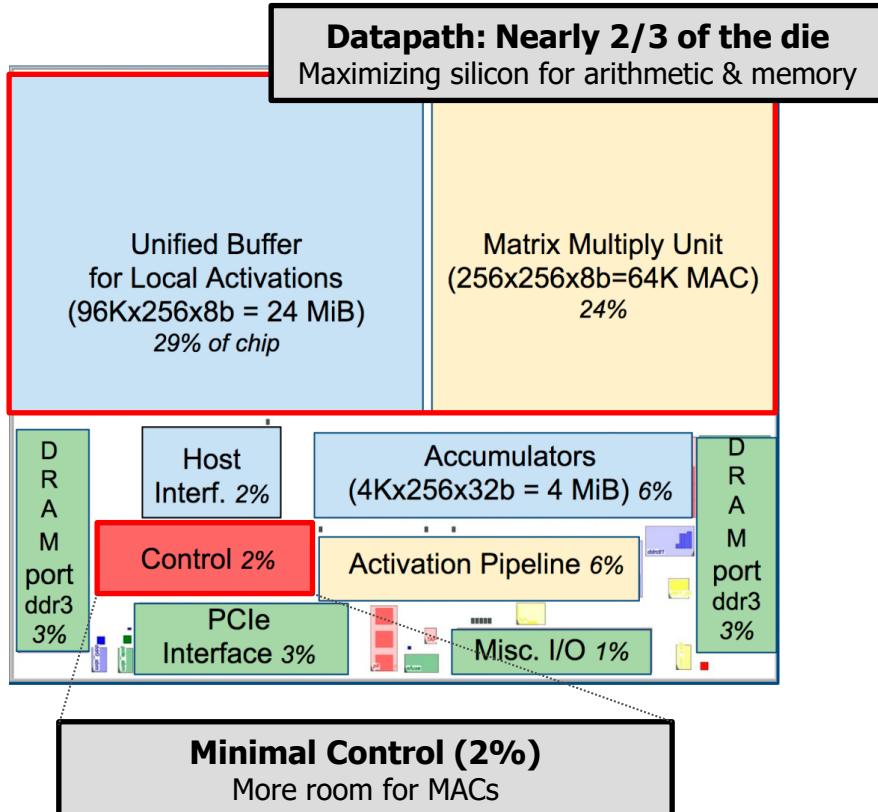
Figure 4. Systolic data flow of the Matrix Multiply Unit.

### MXU: $256 \times 256$ Systolic Array of 8-bit MACs

- **Mechanism:** Activations flows in from the left, Weights from the top (Diagonal Wavefront).
- **Efficiency:** "**Read Once, Use Many**" strategy reduces expensive 'United Buffer' accesses.

# Architecture

## Floorplan of TPU die

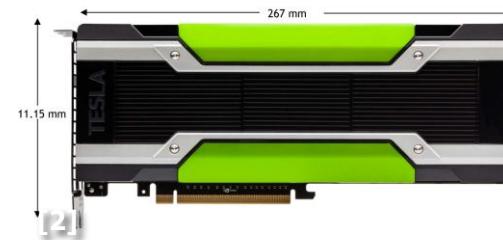


Component	Die Area	Key Role & Characteristics
<b>Unified Buffer</b>	<b>29%</b>	<b>Data Reuse</b>
<b>Matrix Multiply Unit</b>	<b>24%</b>	<b>Compute Power</b>
<b>Control Logic</b>	<b>2%</b>	<b>Minimalism</b> <ul style="list-style-type: none"> <li>• No caches or branch prediction</li> <li>• Saves space for more arithmetic units</li> </ul>

# **Methodology**

# Methodology

Comparing TPU with contemporary datacenter hardware



Metric	CPU (Baseline)	GPU (Accelerator)	TPU (Custom ASIC)
<b>Product</b>	Intel Haswell E5-2699 v3	Nvidia K80 (Tesla)	Google TPU v1
<b>Technology</b>	22 nm	28 nm	28 nm
<b>Die Size</b>	662 mm <sup>2</sup>	561 mm <sup>2</sup> (per die)	< 331 mm <sup>2</sup>
<b>TDP (Power)</b>	145 W	150 W (per die)	<b>75 W</b>
<b>Memory</b>	51 MiB Cache	16 MiB + 12 GB GDDR5	28 MiB + 8 GB DDR3
<b>Peak Ops/s</b>	2.6 TOPS (8b)	2.8 TOPS (FP)	<b>92 TOPS (8b)</b>

[1] Intel, "Intel Haswell E5-2699 v3"

[2] NVIDIA, "Tesla K80 GPU Accelerator Board Specification"

# Methodology

## Datacenter workloads

엄격한 응답시간  
제한이라는 조건 하라는  
걸 먼저 소개

엄격한 응답시간  
제한이라는 조건 하에서  
tpu가 세팅한 결과값

Constrained by **7ms latency limit**

App Name	% of Deployed TPUs in July 2016	Total Layers	TPU Ops / Weight Byte	TPU Batch Size
MLP0	61%	5	200	200
MLP1		4	168	168
LSTM0	29%	24	64	64
LSTM1		37	96	96
CNN0	5%	16	2,888	8
CNN1		89	1,750	32

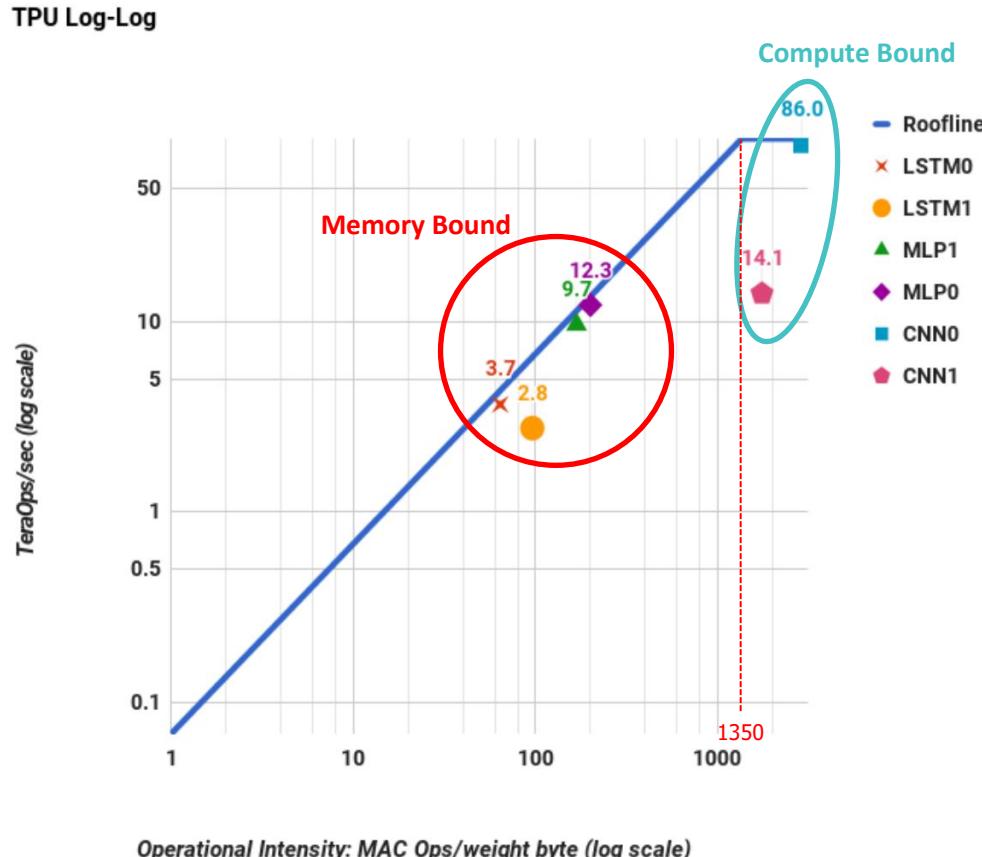
Type	Batch	99th% Response	Inf/s (IPS)	% Max IPS
CPU	16	7.2 ms	5,482	42%
CPU	64	21.3 ms	13,194	100%
GPU	16	6.7 ms	13,461	37%
GPU	64	8.3 ms	36,465	100%
TPU	200	7.0 ms	225,000	80%
TPU	250	10.0 ms	280,000	100%

Table 4: MLP0 Performance: Throughput vs. Latency Constraints

# Evaluation

# Evaluation

## TPU Performance analysis: The roofline model



Region	Memory Bound	Compute Bound
Criteria (Intensity)	Intensity < 1,350 Ops/Byte	Intensity > 1,350 Ops/Byte
Models	LSTM0, LSTM1, MLP0, MLP1	CNN0, CNN1
Metric	MLP0	LSTM1
TeraOps/sec	12.3	2.8
Useful MACs	12.5%	6.3%
Weight Stall	53.9%	62.1%
	CNN0 (Inception)	CNN1 (AlphaGO)
	86.0	14.1
	78.2%	22.5%
	0.0%	28.1%

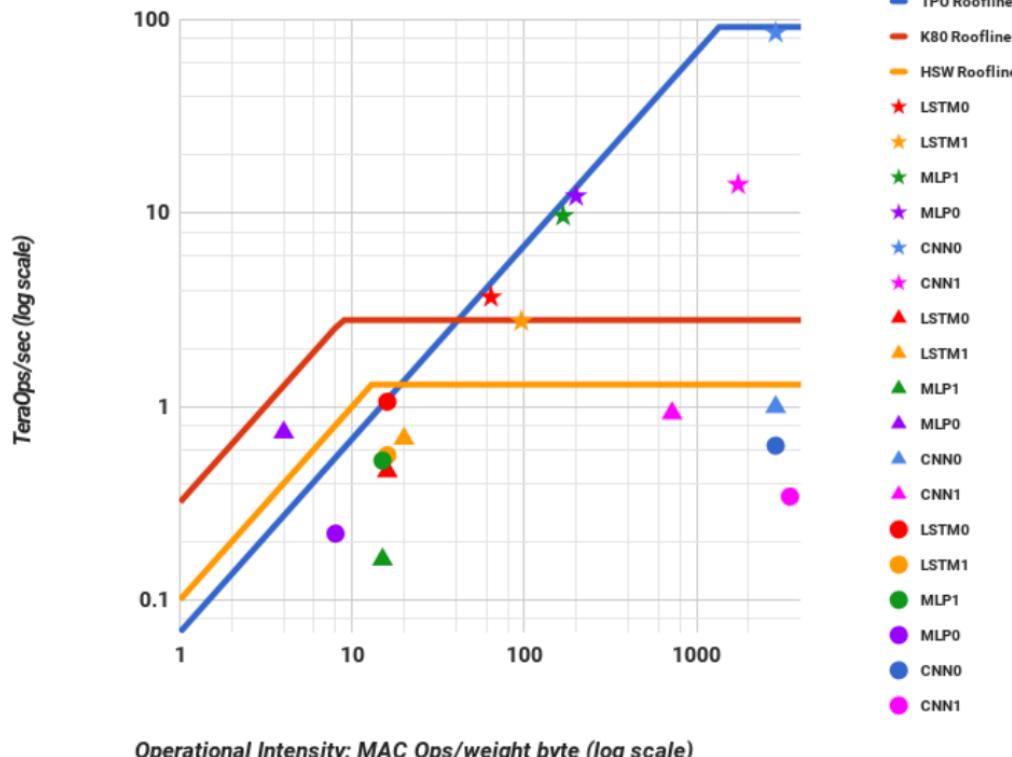
### Performance Gap Analysis (CNN0 vs CNN1)

- Vertical Depth: CNN1 (89 layers) >> CNN0 (16 layers)
- Horizontal Depth: CNN1 is "Shallow" → Low MMU Utilization

# Evaluation

## Roofline comparison (TPU vs. Others)

Log-Log Scale

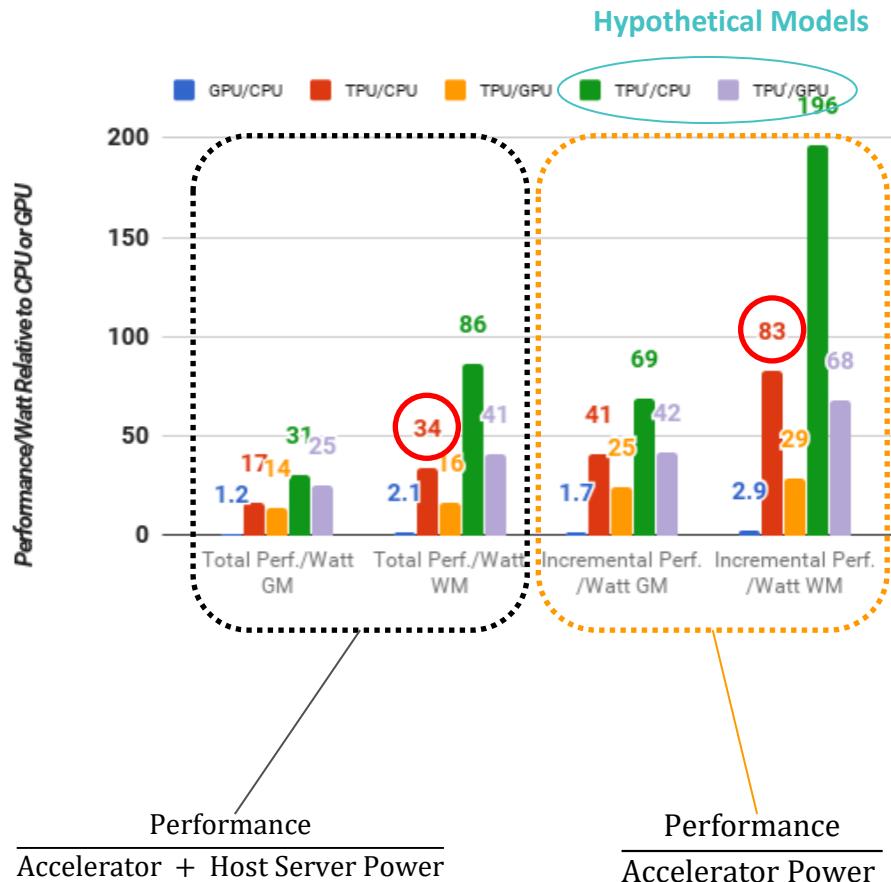


Deployed TPUs	61%		29%		5%		GM	WM
Type	DNN	LSTM	CNN	DNN	LSTM	CNN		
GPU	2.5	0.3	0.4	1.2	1.6	2.7	1.1	1.9
TPU	41.0	18.5	3.5	1.2	40.3	71.0	14.5	29.2
Ratio	16.7	60.0	8.0	1.0	25.4	26.3	13.2	15.3

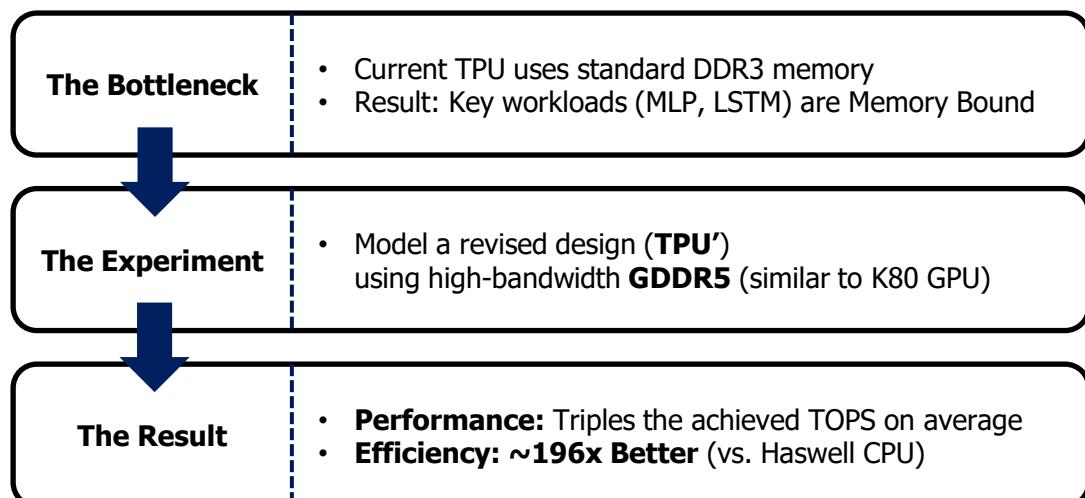
Feature	Geometric Mean (GM)	Weighted Mean (WM)
Definition	Equal weight (1/6) per app.	Weighted by actual usage (%).
Key Driver	Rare apps (CNN) count equally.	Top apps (MLP0 61%) dominate.
TPU Speedup (vs. CPU)	<b>14.5x.</b>	<b>29.2x.</b>
Conclusion	Undervalues performance.	Shows <b>real-world impact.</b>

# Evaluation

## Relative TDP of each server (CPU vs. GPU vs. TPU)



Feature	Total Perf/Watt	Incremental Perf/Watt
<b>Formula</b>	Accelerator + Host Server Power	Accelerator Power <b>Only</b>
<b>Focus</b>	Real datacenter electric bill (TCO).	True efficiency of the TPU chip.
<b>TPU Impact (WM)</b>	<b>~30x</b> Better (vs. Haswell CPU)	<b>~83x</b> Better (Diluted by host power). <b>(Justifies Custom ASIC)</b>



# Evaluation

## Energy proportionality of workload utilization



### The Observation

- CPU (Haswell):** Power drops significantly when idle (56% power at 10% load).
- TPU:** Power stays high even when idle (**88% power** at 10% load)

### The Reason

- Short development schedule (**15 months**) prevented including complex energy-saving features.

### The Implication

- TPU – “Poor Energy Proportionality”**
- To maximize ROI, the TPU workload queues must be kept full (100% utilization).

# Evaluation

## Performance sensitivity & Prospective of TPU

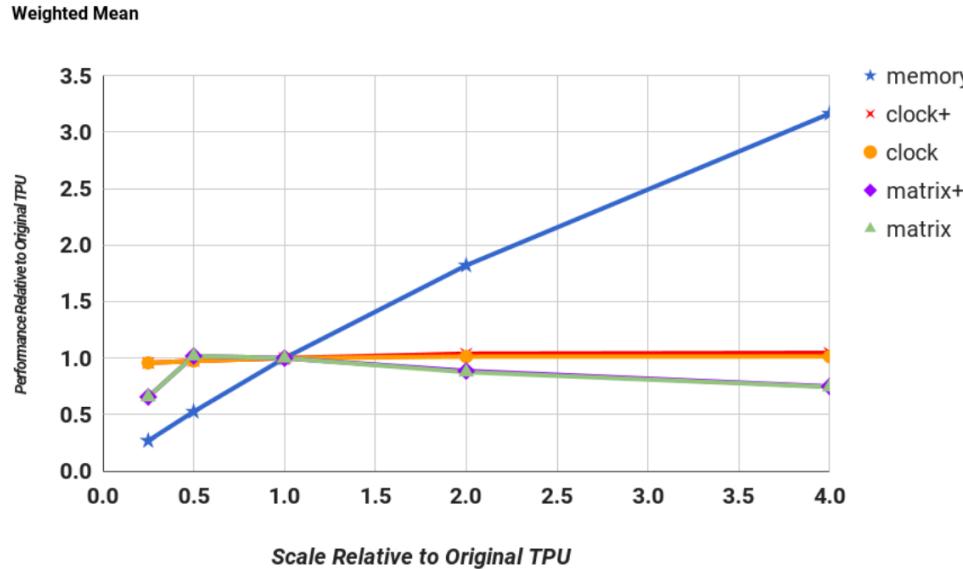
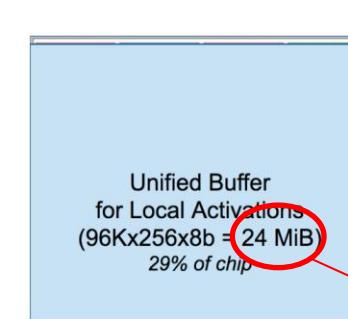


Figure 11. Weighted mean TPU performance as metrics scale

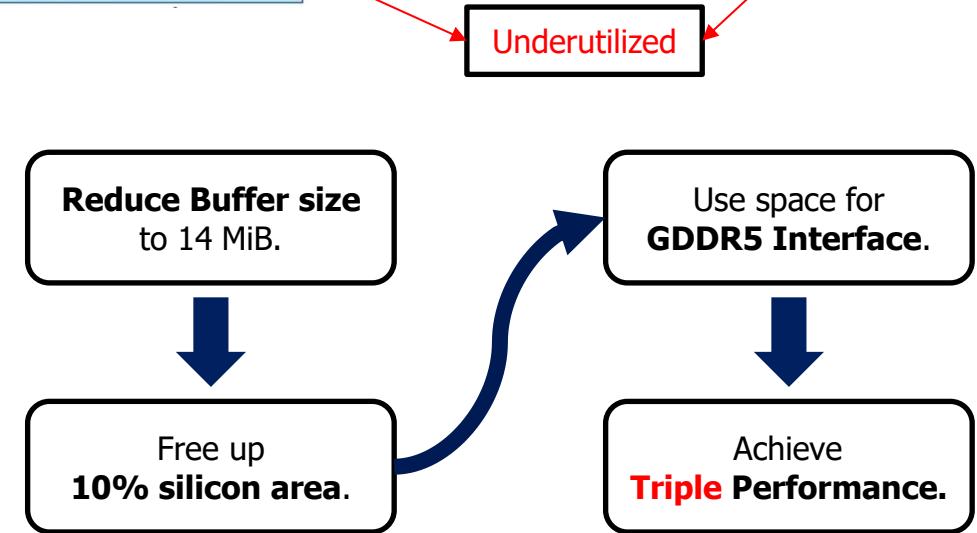


"Simulation Error rate: < 8%"



MLP0	MLP1	LSTM0	LSTM1	CNN0	CNN1
11.0	2.3	4.8	4.5	1.5	13.9

Table 8. Maximum MiB of the 24 MiB Unified Buffer used per NN application. A 14 MiB Unified Buffer is sufficient for them now, due to improvements in the new software allocator.



# Conclusion

# Conclusion

## Summary and Key Takeaways

Category	Key Factors	Metrics / Evidence
<b>1. Quantitative Impact</b>	<b>Speedup</b>	<ul style="list-style-type: none"> <li>vs CPU/GPU: <b>Avg 15x ~ 30x</b> Faster</li> </ul>
	<b>Efficiency</b>	<ul style="list-style-type: none"> <li>TOPS/Watt: <b>30x ~ 80x</b> Higher</li> </ul>
	<b>Cost / TCO</b>	<ul style="list-style-type: none"> <li><b>Order-of-Magnitude</b> Better Cost-Performance</li> </ul>
<b>2. Success Factors</b>	<b>Minimalism</b>	<ul style="list-style-type: none"> <li>No Caches, No Branch Predict, No Out-of-Order</li> </ul>
	<b>Massive Compute</b>	<ul style="list-style-type: none"> <li><b>65,536 (64K)</b> 8-bit MAC Matrix Unit</li> </ul>
	<b>On-Chip Memory</b>	<ul style="list-style-type: none"> <li><b>28 MiB</b> Software-Managed Memory</li> </ul>
<b>3. Key Point</b>	<b>Determinism</b>	<ul style="list-style-type: none"> <li>Meets <b>99th-percentile Latency</b> Target</li> </ul>
	<b>Memory Bound</b>	<ul style="list-style-type: none"> <li>4 of 6 NN Workloads (MLP, LSTM) are <b>Memory-Bound</b></li> </ul>
	<b>TPU' Potential</b>	<ul style="list-style-type: none"> <li><b>3x Performance Boost</b> with <b>GDDR5</b> Memory</li> </ul>